

# Efficient CPU↔GPU data transfers

## CUDA 6.0 Unified Virtual Memory


Juraj Kardoš

(University of Lugano)

July 9, 2014

# Motivation

- Impact of **data transfers** on overall application performance



# When GPU↔CPU memory transfers are performed?

# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**

# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**

Where else?

# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**

Where else?

- Loading **kernel binary code** (implicitly, by driver)

# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**

Where else?

- Loading **kernel binary code** (implicitly, by driver)
- Loading **kernel arguments** (transferred into GPU constant memory upon kernel launch, implicitly, by driver)

# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**

Where else?

- Loading **kernel binary code** (implicitly, by driver)
- Loading **kernel arguments** (transferred into GPU constant memory upon kernel launch, implicitly, by driver)
- Passing return scalar value, e.g. reduction result (remember `__global__` functions are always `void`)



# When GPU↔CPU memory transfers are performed?

- When transferring **input/output arrays**


Where else?

- Loading **kernel binary code** (implicitly, by driver)
- Loading **kernel arguments** (transferred into GPU constant memory upon kernel launch, implicitly, by driver)
- Passing return scalar value, e.g. reduction result (remember `__global__` functions are always `void`)
- Initializing `__device__` variables



# PCI Express overview

- Computer expansion bus
- Point-to-point connection
- Lane sharing
- Single bus (x1)
  - 500 MB/s per lane (PCI-e v2)
- Multiple lanes (x2, x4, x8, x16, x32)
  - 8 GB/s for a 16 lane bus




# Generations of PCI-Express


PCI Express version	Per Lane Bandwidth	x16 Bandwidth
1.0 (2003)	250 MB/s	4 GB/s
2.0 (2007)	500 MB/s	8 GB/s
3.0 (2010)	984 MB/s	15 GB/s
4.0 (2014-15)	1969 MB/s	31 GB/s



# PCI-E Bandwidth Test



# Remember PCI-E Lanes?





# Types of data transfers in CUDA

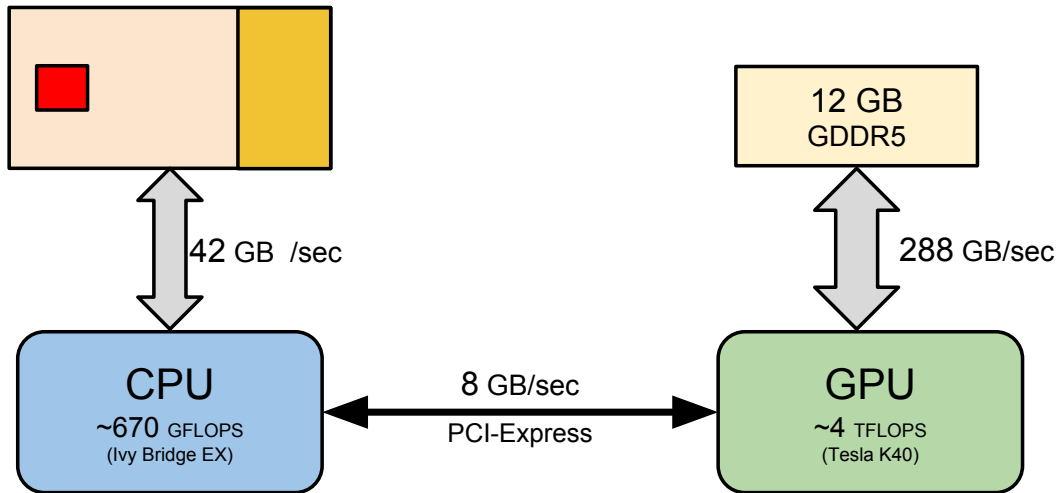
- Pageable or pinned
- Explicit or implicit (automatic, UVM)
- Synchronous or asynchronous
- Peer to peer (between GPUs of the same host)
- GPUDirect (between GPU and network interface)



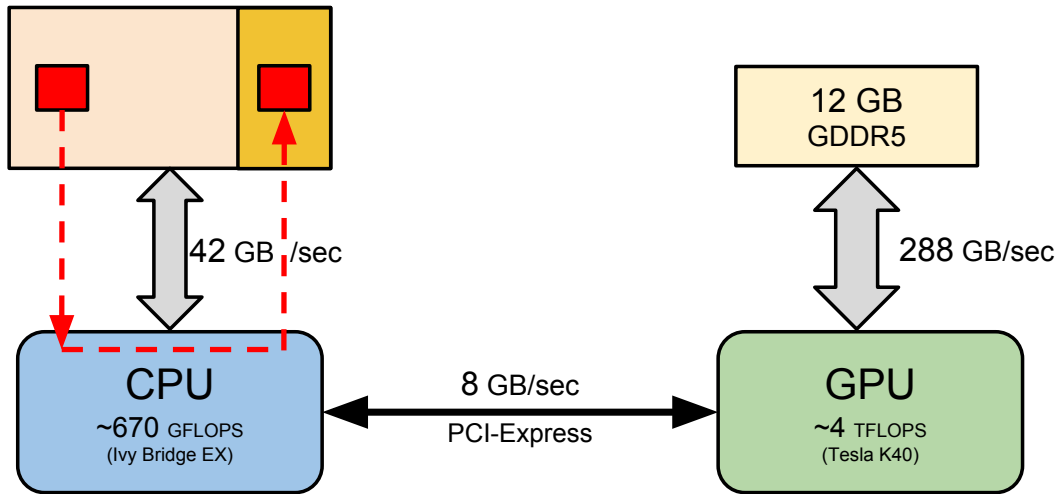
# Types of data transfers in CUDA

- **Pageable or pinned**
- Explicit or implicit (automatic, UVM)
- Synchronous or asynchronous
- Peer to peer (between GPUs of the same host)
- GPUDirect (between GPU and network interface)


# Pageable and pinned memory transfer



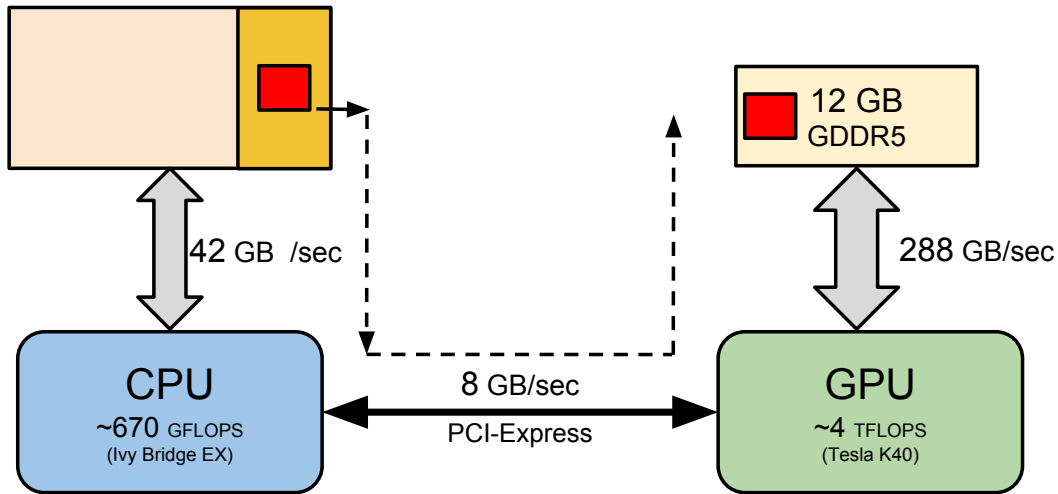
# Pageable and pinned memory transfer




# Pageable and pinned memory transfer



## Pageable and **pinned** memory transfer



# Pageable and pinned memory transfer



# Pageable and pinned memory transfer

```
//allocate memory
w0 = (real*)malloc( szarrayb);
cudaMalloc(&w0_dev, szarrayb);

//memcpy
cudaMemcpy(w0_dev, w0, szarrayb, ←
           cudaMemcpyHostToDevice);

//kernel compute
wave13pt_d<<<...>>>( ..., w0_dev, ...);

//memcpy
cudaMemcpy(w0, w0_dev, szarrayb, ←
           cudaMemcpyDeviceToHost);
```

Listing 1: Pageable

```
//allocate memory
cudaMallocHost(&w0, szarrayb);
cudaMalloc(&w0_dev, szarrayb);

//memcpy
cudaMemcpy(w0_dev, w0, szarrayb, ←
           cudaMemcpyHostToDevice);

//kernel compute
wave13pt_d<<<...>>>( ..., w0_dev, ...);

//memcpy
cudaMemcpy(w0, w0_dev, szarrayb, ←
           cudaMemcpyDeviceToHost);
```

Listing 2: Pinned

# Pageable and pinned memory transfer - Summary

- Pageable memory - user memory space, requires extra mem-copy
- Pinned memory - kernel memory space
- Pinned memory performs better (higher bandwidth)
- Do not over-allocate pinned memory - reduces amount of physical memory available for OS




# Types of data transfers in CUDA

- Pageable or pinned
- **Explicit or implicit (UVM)**
- Synchronous or asynchronous
- Peer to peer (between GPUs of the same host)
- GPUDirect (between GPU and network interface)

# Unified Memory

- Developer view on memory model
- Still two distinct physical memories on HW level



# Unified Memory - Usage

```
//allocate memory
w0 = (real*)malloc( szarrayb);
cudaMalloc(&w0_dev, szarrayb);

//memcpy
cudaMemcpy(w0_dev, w0, szarrayb, ←
    cudaMemcpyHostToDevice);

//kernel compute
wave13pt_d<<<...>>>( ..., w0_dev, ...);

//memcpy
cudaMemcpy(w0, w0_dev, szarrayb, ←
    cudaMemcpyDeviceToHost);

//host function
f(w0);
```

Listing 3: Explicit memory

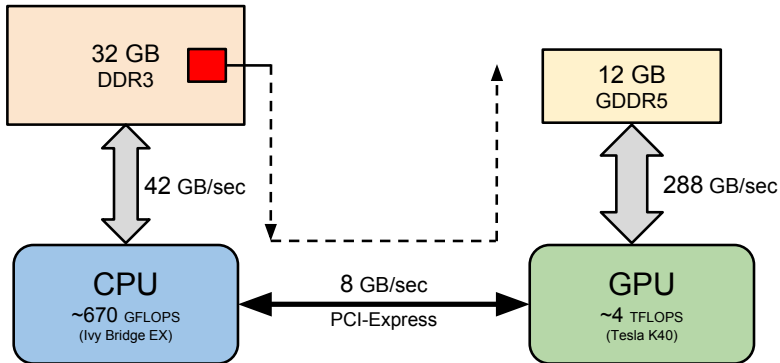
```
//allocate memory
cudaMallocManaged(&w0, szarrayb);

//kernel compute
wave13pt_d<<<...>>>( ..., w0, ...);


//host function
f(w0);
```

Listing 4: UVM


# Unified Memory - Use Case



# Unified Memory - Use Case




# Unified Memory - Use Case

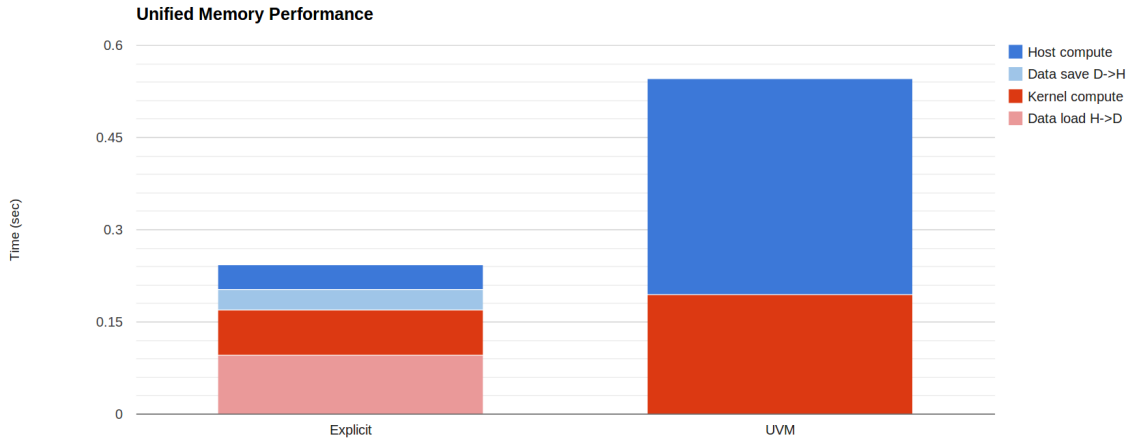


# Unified Memory - Use Case

- How does UVM perform when compared to explicit memory movements?



# Implicit memory transfers: UVM






# Implicit memory transfers: UVM

- How does UVM perform in case of multi-threading?

# Implicit memory transfers: UVM

- UVM Implements CS - threads are serialized, performance degradation



# UVM - Summary

---

- Simplifies programming model, but...

# UVM - Summary


---

- Simplifies programming model, but...
- Performance issue D  $\rightarrow$  H
- CS in multi-threaded application

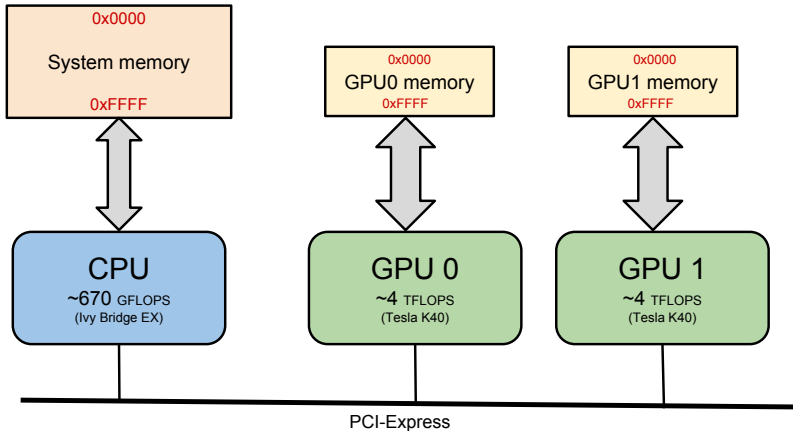
# Types of data transfers in CUDA

- Pageable or pinned
- Explicit or implicit (automatic, UVM)
- Synchronous or asynchronous
- **Peer to peer (between GPUs of the same host)**
- GPUDirect (between GPU and network interface)

# Peer to peer data transfers overview

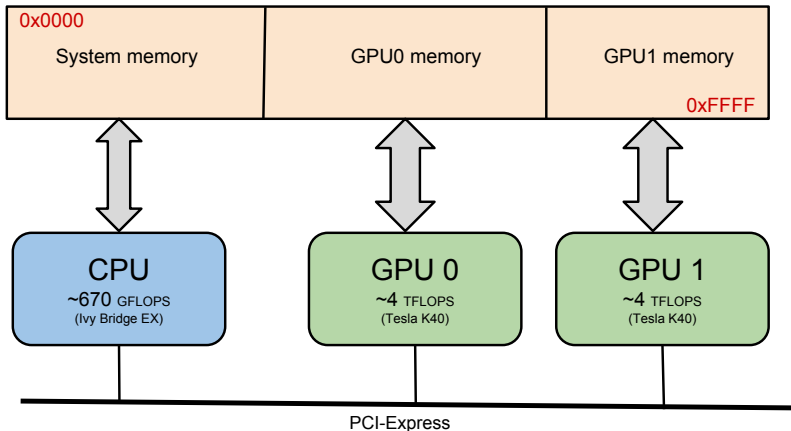


# Peer to peer data transfers - Unified Virtual Addressing



# Peer to peer data transfers - Unified Virtual Addressing

- UVA maps memories into single address space





# P2P Memory Transfer - Usage

```
//allocate memory on gpu0 and gpu1
cudaSetDevice(gpuid_0);
cudaMalloc(&gpu0_buf, buf_size);
cudaSetDevice(gpuid_1);
cudaMalloc(&gpu1_buf, buf_size);

//enable P2P
cudaSetDevice(gpuid_0);
cudaDeviceEnablePeerAccess(gpuid_1, 0);
cudaSetDevice(gpuid_1);
cudaDeviceEnablePeerAccess(gpuid_0, 0);

//P2P copy
cudaMemcpy(gpu0_buf, gpu1_buf, buf_size, cudaMemcpyDefault)
```

Listing 5: P2P

# Peer to peer data transfers - Summary

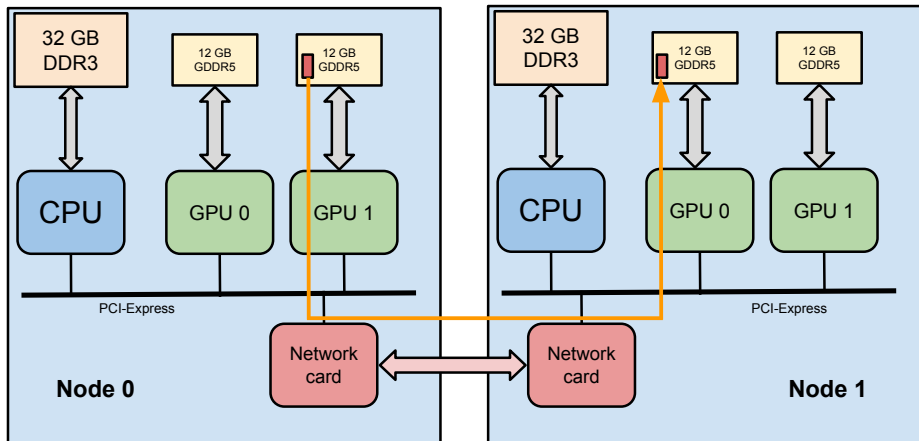
- P2P and UVA can be used to both **simplify** and **accelerate** CUDA programs
- One address space for all CPU and GPU memory
  - Determine physical memory location from pointer value
  - Simplified library interface – `cudaMemcpy()`
- Faster memory copies between GPUs with less host overhead

# Types of data transfers in CUDA

- Pageable or pinned
- Explicit or implicit (automatic, UVM)
- Synchronous or asynchronous
- Peer to peer (between GPUs of the same host)
- **GPUDirect (between GPU and network interface)**

# GPU direct overview

- Eliminate CPU bandwidth and latency bottlenecks using remote direct memory access transfers between GPUs and other PCIe devices



# General recommendations

- PCI-E is efficient only starting from reasonably large data buffer
- UVM simplifies programming model but may result in worse performance
- It's always a good idea to know when underlying runtime routes data through intermediate buffer (additional copying) and avoid that (pinned memory, GPUDirect)
- It's always a good idea to compute something, while data is being transferred (asynchronous)

# Control questions

- 1 How many PCI-E lanes 1 GPU can consume? Suppose you have 40 PCI-E lanes and 4 GPUs. How many lanes there will be available per GPU, if they all are transferring data simultaneously?
- 2 Given that UVM is slower than explicit copying, what it could still be good for?
- 3 What is better to use for multi-gpu application: P2P memory transfers, GPUDirect or CUDA-aware MPI?

## Control questions: answers

- 1** 1 GPU usually can use up to  $16\times$  lanes. With 4 GPUs in a single system, there will be  $8\times$  lanes link per GPU in average, i.e. 2 times less than with single GPU in system. Note this when building your GPU servers.
- 2** UVM simplifies GPU porting, allowing you omit explicit memory copies during intensive GPU kernels code development.
- 3** CUDA-aware MPI uses P2P and GPUDirect as underlying engines. Thus, CUDA-aware MPI might better suite MPI applications, while single-node programs could be written in simpler way with CUDA P2P.