

Notes from learning CUDA by example of wave equation stencil

Dmitry Mikushin, Pietro Benedusi

(University of Lugano)

July 9, 2014

4th assignment of PDCLab course @ USI

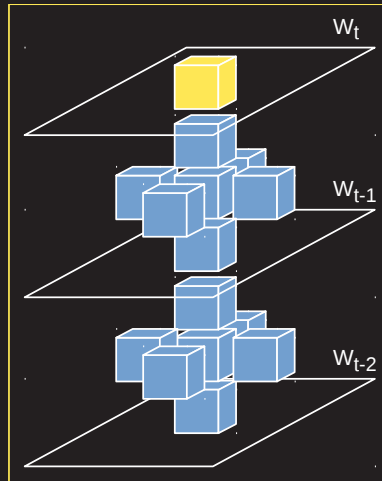
Pick up your favourite test program from previous assignments (Game Of Life (2d), wave13pt (3d), or hidden parallelism kernel (2d)) and port it to:

- CUDA
- Cell Broadband Engine
- **(optional)** Xeon Phi

4th assignment of PDCLab course @ USI

- Time iterations (“forget” the oldest data and replace it with newest)
- 3D Loop on every iteration:

```
for (int k = 2; k < ns - 2; k++)  
  for (int j = 2; j < ny - 2; j++)  
    for (int i = 2; i < nx - 2; i++)  
    {  
      w2[k][j][i] = m0 * w1[k][j][i] - w0[k][j][i] +  
  
      m1 * (  
        w1[k][j][i+1] + w1[k][j][i-1] +  
        w1[k][j+1][i] + w1[k][j-1][i] +  
        w1[k+1][j][i] + w1[k-1][j][i]) +  
  
      m2 * (  
        w1[k][j][i+2] + w1[k][j][i-2] +  
        w1[k][j+2][i] + w1[k][j-2][i] +  
        w1[k+2][j][i] + w1[k-2][j][i]);  
    }
```



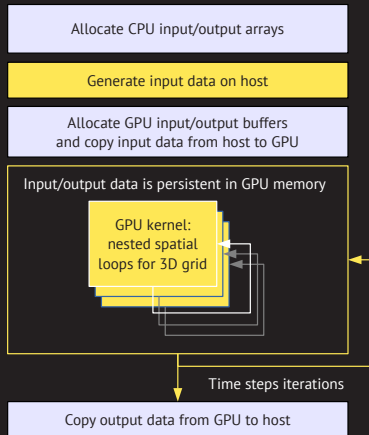
Notes on understanding the algorithm

Familiarity with data flow is crucial:

- What are the inputs and the outputs
- How propagation moves from current level to the next one

Ways to explore the program behavior:

- Run the code in debugger line by line



Notes on understanding CUDA

1 Study simple examples, e.g. `cuda_gpu_test` and `compute_grid_3d`

2 How to select the compute grid config?

- Choose grid point : GPU thread mapping: 1:1 or N:1 (with jumping by `blockDim.x` increment)
- Use blocks of size divisible by `warpSize`, better to have more smaller blocks, than less larger

3 How to handle multiple dimensions?

- Either linearize into 1D and recover (i,j,k) from 1D index in GPU kernel:

```
wave13pt<<<szarray/THREADS_PER_BLOCK,THREADS_PER_BLOCK>>>(  
    nx, ny, ns, m0, m1, m2, d_w0, d_w1, d_w2);
```

- Or use multidimensional grid, that allows to map multidimensional problems naturally:

```
wave13pt<<< X, Y, Z >>>(...)
```

Notes on bug hunting: basic checks

I've got the working version with wrong results. Where should I search for an error?

- 1 Check error status of every CUDA function call with e.g. `CUDA_ERR_CHECK` macro
- 2 Check error status of every CUDA kernel with `CUDA_ERR_CHECK(cudaGetLastError())`
- 3 Try to run only single time step
- 4 Use `printf()` to check if swapping of time levels is done correctly
- 5 You can also do `printf()` from GPU

Notes on bug hunting: synchronization

- Check if the CUDA kernel is synchronized properly:

```
for (int it = 0; it < nt; it++)  
{  
    wave13pt<<<config.gridDim, config.blockDim>>>(  
        nx, ny, ns,  
        config,  
        m0, m1, m2, w0p, w1p, w2p);  
    CUDA_SAFE_CALL(cudaGetLastError());  
    CUDA_SAFE_CALL(cudaDeviceSynchronize());  
  
    real* w = w0p; w0p = w1p; w1p = w2p; w2p = w;  
    int idx = idxs[0]; idxs[0] = idxs[1]; idxs[1] = idxs[2]; idxs[2] = idx;  
}
```

- Pointer swapping does not involve CUDA functions \Rightarrow dependency will not be detected \Rightarrow incorrect result without explicit synchronization.

Notes on bug hunting: cuda-memcheck

- Detect out-of-range memory accesses with cuda-memcheck \approx valgrind for GPU

```
$ cuda-memcheck ./wave13pt 512 256 256 10
===== CUDA-MEMCHECK
m0 = 0.680375, m1 = -0.035206, m2 = 0.094366
initial mean = 0.000024
init time = 0.219432 sec
device buffer alloc time = 0.023246 sec
data load time = 0.062911 sec (5.960818 GB/sec)
===== Invalid __global__ read of size 4
=====      at 0x000012f0 in /home/marcusmae/apc/trainings/schools/↵
      cscs_summer_school_2014/slides/pietro/black/wave13pt_gpu/cuda/../../wave13pt.cu:642:↵
      wave13pt
=====      by thread (95,0,0) in block (2,237,251)
=====      Address 0xb181409bc is out of bounds
```


Notes on performance

- 1 CUDA code for GPU is compiled with full optimization by default (i.e. `implicit -O3`)
- 2 Initially the perf was poor, until the block size of `{128, 1, 1}` was selected
- 3 Always maintain memory coalescing