



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

MPI: Non-blocking point-to-point communication

CSCS-USI Summer School 2014

Themis Athanassiadou

Serpiano, Switzerland

Imagine the following situation...

- Each process sends a message to the right neighbor, and receives from the left... (live demo!)
- The code for each process is: `MPI_Send(...., to_right,...)`
`MPI_Recv(...., from_left...)`




Image: Rolf Rabenseifner, HLRS

Imagine the following situation...

- Each process sends a message to the right neighbor, and receives from the left... (live demo!)
- The code for each process is: `MPI_Send(..., to_right,...)`
`MPI_Recv(..., from_left...)`




Image: Rolf Rabenseifner, HLRS






Non-blocking communication to the rescue...


Non-blocking mode (**MPI_Isend**, **MPI_Irecv**):

- Non-blocking send and receive routines behave similarly - they will return **I**mmediately. They do not wait for any communication events to complete
- Non-blocking operations simply "request" the MPI library to perform the operation when it can. The user cannot predict when this will happen.
- It is unsafe to modify the application buffer until communication is actually completed. Completion is ensured by the wait functions.
- Non-blocking communications are primarily used to overlap computation with communication.

The ring example: Avoiding the deadlock

The code now for each process is:

MPI_Isend(...., to_right,...) 
MPI_Recv(...., from_left...) 
MPI_Wait (.....) 





MPI_Isend/ MPI_Irecv (FORTRAN version)

call MPI_ISEND(buf, count, type, dest, tag, comm, req, ierr)
call MPI_IRECV(buf, count, type, source, tag, comm, req, ierr)

- **buf** : array
- **count** (INTEGER) number of element of buf to be sent
- **type** MPI type of buf (MPI_INTEGER, MPI_REAL etc)
- **dest** (INTEGER) rank of the destination process
- **tag** (INTEGER) number identifying the message
- **comm** (INTEGER) communicator of the sender and receiver
- **req** (INTEGER) **output**, identifier of the communications handle
- **ierr** (INTEGER) **output**, error code (if ierr=0 no error occurs)



MPI_Isend, MPI_Irecv (C Version)

```
int MPI_Isend(void *buf, int count, MPI_Datatype  
             type, int dest, int tag, MPI_Comm comm, MPI_Request *req);
```

```
int MPI_Irecv (void *buf, int count, MPI_Datatype  
              type, int source, int tag, MPI_Comm comm, MPI_Request *req);
```



Waiting for completion

- **Fortran:**

`MPI_WAIT(req, status, ierr)`

`MPI_WAITALL (count,array_of_requests,array_of_statuses, ierr)`

- A call to this subroutine causes the code to wait until the communication referred to by **req** is complete.
- `req(INTEGER)`: input/output, identifier associated to a communications event (initiated by `MPI_ISEND` or `MPI_IRECV`).
- `Status(INTEGER)` array of size `MPI_STATUS_SIZE`, if `req` was associated to a call to `MPI_IRECV`, `status` contains information on the received message, otherwise `status` could contain an error code.
- `ierr(INTEGER)` output, error code (if `ierr=0` no error occurs).

- **C:**

`int MPI_Wait(MPI_Request *req, MPI_Status *status)`

`int MPI_Waitall (count,&array_of_requests,&array_of_statuses)`



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

Thank you for your attention.