



CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Advanced topics in OpenACC

CSCS-USI Summer School 2014

Markus Wetzstein
wetzstein@cscs.ch



Motivation and Outlook

there surely has to be more to OpenACC...???

...yes, indeed:

- **tuning workshare constructs**
- **reduction operations**
- **asynchronous operations**
- **interoperability with CUDA and libraries**

I have an existing OpenMP code, how easy is it to port to OpenACC?

...fairly easy, BUT:

**important differences between OpenMP and OpenACC
(although it all looks very similar)**



Parallel region with multiple workshared loops ?

- **often used construct in OpenMP**
- **thread creation/ destruction @ begin/end of parallel region**
- **loops are workshared**
- **code between loops is executed redundantly on each thread**



what about OpenACC ?

```
!$omp parallel shared(...) &  
!$omp& private(...) [...]  
  
!$omp do  
do i=istart,iend  
    ...  
enddo  
[some code, not workshared]  
!$omp do  
do j=jstart,jend  
    do k=kstart,kend  
        ...  
    enddo  
enddo  
!$omp end parallel
```



Use of parallel and loop

- **OpenACC parallel region can span several loops**
BUT
- **OpenACC \neq OpenMP although visually very similar**
- **each loop executed by the same sets of threads**
- **no synchronization between these and none with code in between !**
- **no global barrier mechanism inside parallel region**
- **very easy to create race conditions**

```
!$acc parallel shared(...) &  
!$acc& private(...) [...]  
  
!$acc loop  
do i=istart,iend  
    ...  
enddo  
[some code, not workshared]  
!$acc loop  
do j=jstart,jend  
    do k=kstart,kend  
        ...  
    enddo  
enddo  
!$acc end parallel
```



Use of parallel and loop (2)

```
!$acc parallel shared(...)
private(...) [...]

!$acc loop
do i=istart,iend
    ...
enddo

[some code, not workshared]
!$acc loop
do j=jstart,jend
    do k=kstart,kend
        ...
    enddo
enddo

!$acc end parallel
```



```
!$omp parallel shared(...) &
!$omp private(...) [...]

!$omp do
do i=istart,iend
    ...
enddo
!$omp end do nowait

[some code, not workshared]
!$omp do
do j=jstart,jend
    do k=kstart,kend
        ...
    enddo
enddo
!$omp end do nowait
!$omp end parallel
```

Use of parallel and loop: recommendations

- **only use composite `parallel` loop at first**
→ get correct code
- **understand `loop` by itself as corresponding to having an implicit `nowait` clause in OpenMP (without the option to have a `wait`!)**
- **carefully separate directives as a later performance tuning step:**
- **only if** you're sure the loops can be **independent** kernels and have no race conditions !
- consider explicitly using `async` clause instead
→ slightly more complicated, but enhanced code clarity

```
!$acc parallel shared(...) &  
!$acc& private(...) [...]  
  
!$acc loop  
do i=istart,iend  
    a(i)=b(i)*const + c(i)  
enddo  
[some code, not workshared]  
!$acc loop  
do j=jstart,jend  
    do k=kstart,kend  
        x(j,k)=y(j,k)+z(j,k)  
    enddo  
enddo  
!$acc end parallel
```



- ✓ loops independent!
- ✓ potentially slightly faster than having two separate `parallel` `loop` clauses



Kernels vs parallel regions

- **what is the difference between `kernels` and `parallel` ?**
- very similar usage
- have different historic origins (`kernels` → PGI, `parallel` → Cray)
- OpenACC standard not very helpful to understand when to use what

```
#pragma acc parallel loop [...]  
for(i=istart ;i<=iend; i++) {  
    ...  
}  
#pragma acc kernels loop [...]  
for(i=istart ;i<=iend; i++) {  
    ...  
}
```

in common:

- **both define a region to be accelerated**

differences:

- **different levels of obligation to compiler**

<code>parallel</code>	<code>kernels</code>
1 kernel	≥ 1 kernel(s)
must be accelerated	can be accelerated
tuning clauses	no tuning clauses



Kernels vs parallel regions (2)

- **compiler will automatically analyze all loops inside `kernels`**
- **BUT: with kernels, first loop is guaranteed to have finished before second starts**

```
#pragma acc kernels [...] {  
  for(i=istart ;i<iend; i++) {  
    a[i]=b[i]*c[i];  
  }  
  for(i=istart+1 ;i<iend-1; i++) {  
    d[i]= 0.5*(a[i-1]+a[i+1]);  
  }  
} //acc kernels
```

What to use...?

- **`parallel` offers greater control**
- **`kernels` maybe better to initially explore parallelism**

suggestion:

- **don't mix them unless you're really aware of the subtle differences**



Tuning parallel execution

```
[parallel] loop [gang] [worker] [vector]
```

- **parallel execution structured into hierarchy:**
gang → **worker** → **vector**
- **code is executed in parallel with current level of parallelism until a new level is opened**
(**gang**: redundant execution)
- **optional, compilers define defaults (possibly using heuristics)**
- **only allowed on **loop** directive**

```
!$acc parallel loop [...] gang
do i=istart,iend
!$acc loop worker
    do j=jstart,jend
!$acc loop vector
        do k=kstart,kend
            ...
        enddo
    enddo
enddo
!$acc end parallel
```

OpenACC	CUDA
gang	threadblock
worker	warp of threads
vector	threads



when to use it, and why?

THIS is not a very efficient use case !

Tuning parallel execution (2)

Explicitly using multiple levels of parallelism:

- loop iterations must be data independent (except **reductions**)
- usage: indirect indexing, ...
- **worker** and **vector** loops have an implied barrier at end of loop

```
loop [collapse(Nlevels)]
```

- specifies how many tightly nested loops are associated with a **loop** construct
- without **collapse** a **loop** construct only affects the immediately following loop

```
!$acc parallel loop [...] gang
do i=istart,iend
    inew=index_list(i)
!$acc loop worker
    do j=jstart,jend
        jnew=index_list2(j)
!$acc loop vector
        do k=kstart,kend
            a(k,j,i)=b(k,jnew)+c(k,inew)
        enddo
    enddo
enddo
!$acc end parallel
```

```
!$acc parallel loop [...] collapse(3) &
!$acc& gang worker vector
do i=istart,iend
    do j=jstart,jend
        do k=kstart,kend
            a(k,j,i)=b(k,j)+c(k,i)
        enddo
    enddo
enddo
!$acc end parallel
```



Tuning parallel execution (3)

```
parallel [num_gangs(N1)] [num_workers(N2)] [vector_length(N3)]
```

- **num_gangs**: nr. of gangs to use for parallel region (integer)
- **num_workers**: nr. of workers to use for **worker** loops (integer)
- **vector_length**: nr. of threads to use for **vector** loop (integer)
- **binds to parallel, not loop**
- **if omitted, compiler chooses itself**
- **vector_length**: **compiler might allow only certain values**
e.g. Cray: 1, 64, 128 (default), 256, 512, 1024
- **Cray only allows:**
either num_workers (fixes **vector_length=32**)
or vector_length (fixes **num_workers=vector_length/32**)



Tuning parallel execution (4)

Some suggestions:

- **explicitly using `worker` as separate level often not very useful** (in current implementations)
- **tightly nested loops: try if `collapse` improves performance**
- **tuning `num_gangs` | `num_workers` | `vector_length`**
 - is time consuming
 - optimal choice depends on actual loop
 - focus on expensive loops
- **to debug kernel by running single thread, use:**
`#pragma acc parallel num_gangs(1) vector_length(1)`



Reduction operations

```
[parallel | loop] reduction(operator:variable-list)
```

- **OpenACC reductions very similar to OpenMP**
- **reduction only allowed for scalars**
arrays: rewrite to use temporary scalars inside loop nest for reduction
- **reduction variable is private to each thread**
- **combine result over all threads**
e.g. sum, max, min, logical and
- **careful: reduction over gangs only done at end of parallel construct !**

```
#pragma acc parallel loop  
    reduction(+:t)  
for(i=istart;i<=iend;i++) {  
    t=t + a[i] - b[i];  
}
```

Reduction operations (2)

C / C++		Fortran	
operator	initialization	operator	initialization
+	0	+	0
*	1	*	1
max	least	max	least
min	largest	min	largest
&	~0	iand	all bits on
	0	ior	0
^	0	ieor	0
&&	1	.and.	.true.
	0	.or.	.false.
		.eqv.	.true.
		.neqv.	.false.

initialization automatically by compiler (based on operation)

```
#pragma acc parallel [...] {
[ some code ]
#pragma acc loop reduction(+:t)
           gang vector collapse(2)
for(i=istart; i<=iend; i++) {
    for(j=jstart; j<=jend; j++) {
        t = t + a[j,i] - b[j,i];
    }
}
// reduction of t is INCOMPLETE
// using t here=race condition
} // end acc parallel

// using t here is OK
```

be **careful** with **reductions over gangs** before exiting parallel region



Asynchronous operations

- **GPUs have more than one queue (CUDA: stream) into which operations can be entered**
in hardware: Nvidia Fermi 16, Kepler 32 (with better hardware to overlap those), nr. of logical queues even much higher
- **operations in different queues can be executed concurrently**
- **CPU can continue execution immediately after adding an operation to a queue → no need to wait for completion of actual operation**
- **potential performance gains from:**
 - overlapping data transfer with computation on GPU
 - overlapping data transfer with computation on CPU
 - expose more parallelism to the GPU (e.g. multiple kernels and data transfers at the same time)

List of asynchronous clauses / directives

```
wait[(handle-list)] [async(handle)]
```

- synchronisation **directive**
- *handle*: non-negative integer denoting the queue
- *handle-list*: list of handles, can only be used with **wait**
- **wait**: wait until all asynchronous operations have completed
- **wait**(*handle-list*): wait until all asynchronous operations in the queues specified by *handle-list* have completed
- **wait** **async**(*handle*): enters the synchronisation into the queue *handle*

```
[parallel | kernel | enter data | exit data | update] [async[(handle)]]  
[wait[(handle-list)]]
```

- **async**: enters the operation into a default queue
- **async**(*handle*): enters the operation into the queue *handle*
- **wait**: operation starts after all asynchronous operations have completed
- **wait**(*handle-list*): operation starts after all asynchronous operations in the queues specified by *handle-list* have completed
- combinations possible, e.g. **parallel wait async(1)** enter parallel region into queue 1, but don't execute it until all asynchronous operations have completed



Asynchronous example 1

```
[prepare array a on CPU]
#pragma acc enter data async(1) copyin(a)
[prepare array b on CPU]
#pragma acc enter data async(2) copyin(b)
#pragma acc parallel loop async(1) present(a)
for(i=istart ;i<iend; i++) {
    a[i]= [some computation on GPU]
}
#pragma acc exit data copyout(a) async(1)
#pragma acc parallel loop async(2) present(b)
for(j=jstart ;j<jend; j++) {
    b[j]= [some computation on GPU]
}
#pragma acc exit data copyout(b) async(2)
[some computation on CPU]
#pragma acc wait
[continue to use updated a,b on CPU]
```

- simple example with two arrays
- update of arrays independent of each other
 - copy data to GPU
 - compute on GPU
 - copy back to CPU
- this approach can be generalized, e.g. for slices of a larger array

Asynchronous example 2

```

REAL :: a(Nvec,Nchunks),b(Nvec,Nchunks)
!$acc data create(a,b)
DO j = 1,Nchunks
!$acc update device(a(:,j)) async(j)
!$acc parallel loop async(j)
  DO i = 1,Nvec
    b(i,j) = [function of a(i,j)]
  ENDDO
!$acc update host(b(:,j)) async(j)
ENDDO
!$acc wait
!$acc end data
  
```

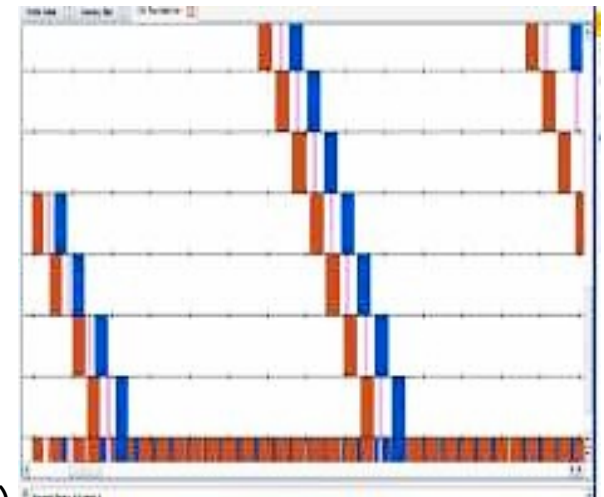
- can overlap 3 streams at once
- use slice number as stream handle
- don't worry if number gets too large
- OpenACC runtime maps it back into allowable range (using MOD function)

Execution times (on Cray XK6):

- CPU: 3.76s
- OpenACC, blocking: 1.10s
- OpenACC, async: 0.34s

NVIDIA Visual profiler:

- time flows left to right
- streams stacked vertically
- only 7 of 16 streams fit in window
- **red**: data transfer to GPU
- **pink**: computational on GPU
- **blue**: data transfer from GPU
- vertical slice shows what is overlapping
- collapsed view at bottom
- async handle modded by number of streams
- so see multiple coloured bars per stream (horizontally)





Recommendations for use of async

- **view it as part of performance tuning**
- **first implement synchronous code, verify it**
- **investigate bottlenecks:**
 - do the kernels need tuning?
 - do the data transfers need tuning?
- **look for data independencies**
 - across kernels
 - between kernels and host code
- **once you have the extent of independent regions, add asynchronous clauses / directives**
- **careful with async handles: only integers and easy to confuse if you need many different ones**
 - consider using e.g. named integer constants if reasonably descriptive naming is possible, e.g. to separate different sets of queues from one another



Use of data on GPU in libraries / CUDA

```
!$acc host_data use_device(var-list)
```

- **how to pass a pointer to memory on the GPU to a library, or to a CUDA kernel? E.g. to:**
 - use third party GPU library (e.g. Cray libsci_acc, cuBLAS, cuFFT, ...) to process data already held on device
 - use optimized CUDA kernel to process data already held on device
 - use optimized MPI library to transfer data across nodes directly between the GPU memories
- **host_data makes a pointer on the device available on the host**
- **nested inside data region which put *var-list* on the GPU**

Interoperability with CUDA

```
PROGRAM main
  INTEGER :: a(N)
  [stuff]
  !$acc data copy(a)
  ! Populate a(:) on device
  ! as before
  !$acc host_data use_device(a)
  CALL dbl_cuda(a)
  !$acc end host_data
  !$acc end data
  [stuff]
END PROGRAM main
```

```
__global__ void dbl_knl(int *c) {
  int i = blockIdx.x*blockDim.x+threadIdx.x;
  if (i < N) c[i] *= 2;
}

extern "C" void dbl_cuda_(int *b_d) {
  cudaThreadSynchronize();
  dbl_knl<<<NBLOCKS,BSIZE>>>(b_d);
  cudaThreadSynchronize();
}
```

- **Call CUDA-C wrapper (compiled with nvcc; linked with normal compiler)**
 - must include cudaThreadSynchronize()
 - Before: so asynchronous accelerator kernels definitely finished
 - After: so CUDA kernel definitely finished before we return to the OpenACC
 - CUDA kernel written as usual
 - Or use same mechanism to call existing CUDA library



Some useful tips at the end...

- **if in doubt, check the OpenACC standard**
- **focus on getting correct code on the GPU first**
- **then start optimizing**
- **focus on data transfers before aiming for a few percent improvement on a kernel**
- **on Cray systems, get detailed info about size of data transfers, kernels launched, etc:**
environment variable `CRAY_ACC_DEBUG=2`
- **make efficient use of tools provided at your computing centre (e.g. DDT/totalview for debugging)**
 - it might take some time to 'learn' using the tool
 - but debugging complex code with `printf` will cost you much more time
 - same goes for performance analysis !